

# DTS 304 – Big Data Computing Project

A Report on: [Integrated Big Data Pipeline for Social Media Sentiment Analysis using Spark, MongoDB, and PostgreSQL](#)

---


Group Name: 12 Nodes of Insight

Member's Student:

ID	Name	Email
30001429	Amanda Gana	amanda.gana@miva.edu.ng
30002640	Christiana Richards	christiana.richards@miva.edu.ng
30000799	Comfort Sophia Iwo Odike	comfort.odike@miva.edu.ng
30010860	King Richard	king.richard@miva.edu.ng
30080232	Aduragbemi Kinoshi	kinoshi.aduragbemi@miva.edu.ng
30012194	Margaret Oliver	Margaret.oliver@miva.edu.ng
30002898	Ojo Ilesanmi	ojo.ilesanmi@miva.edu.ng
30000733	Okon Enang	okon.enang@miva.edu.ng
30025749	Olumide Akinboyewa	olumide.akinboyewa@miva.edu.ng
30062596	Solomon Ayuba	solomon.ayuba@miva.edu.ng
30006647	Susan Ogidan	susan.ogidan@miva.edu.ng
30078164	Yomi Aledare	yomi.aledare@miva.edu.ng

Date of Submission: Wednesday, 13<sup>th</sup> August 2025

Project Repository: [GitHub](#)



# 1. Introduction

This project implements a mini big data pipeline that collects, stores, processes, analyzes, and queries real-time Twitter data. The system integrates three core technologies:

- **MongoDB** for raw tweet storage (NoSQL database)
- **Apache Spark** for large-scale data cleaning and sentiment classification
- **PostgreSQL** for structured storage and relational queries

The chosen theme for data collection was **"technology"** and **"climate"**, both relevant and trending topics.

## 2. System Architecture

The pipeline consists of three major steps:

### 1. Twitter Data Collection & Storage in MongoDB

- Using the Twitter API (via Tweepy), tweets containing "technology" or "climate" are fetched in real-time and stored in MongoDB.
- Raw data is preserved in JSON format for flexibility in downstream processing.

### 2. Data Processing & Sentiment Analysis with Apache Spark

- Tweets are loaded from MongoDB into Spark.
- Text is cleaned by removing URLs, mentions, hashtags, special characters, and extra spaces.
- Sentiment classification is applied using an opinion lexicon (positive/negative word lists).
- Final output is a structured dataset with a *sentiment* column.

### 3. PostgreSQL Integration & Analytical Queries

- The processed dataset is written into a PostgreSQL table.
- SQL queries generate insights such as sentiment distribution, most active users, and time-based sentiment trends.



### 3. Step-by-Step Implementation

#### 3.1 Step 1 – Twitter Data Collection (Tweepy + MongoDB)

**Objective:** To acquire real-time tweets based on keywords and store them for processing.

**Key Actions:**


- Connected to Twitter API v2 using a Bearer Token.
- Search query: *technology OR climate -is:retweet lang:en*

(retrieves English tweets about technology or climate, excluding retweets).

- Stored tweets in a JSON file (*technology\_climate\_tweets.json*) and inserted them into *tweet\_db.tweets* in MongoDB.

**Output:**

- ~80 tweets in MongoDB, each containing:
  - Tweet ID, text, author ID, and creation timestamp.
- Raw dataset available for Spark processing.



#### 3.2 Step 2 - Data Processing & Sentiment Classification (Apache Spark)

**Objective:** To clean tweet text and label each tweet with a sentiment category.

**Key Actions:**

- Loaded data from MongoDB into Spark using the MongoDB Spark Connector.
- Applied regular expression rules to:
  - Remove URLs, mentions, hashtags.
  - Remove special characters and newline breaks
  - Trim leading/trailing spaces.

- Implemented **opinion lexicon-based sentiment analysis**:
  - Used predefined positive and negative word lists.
  - Assigned sentiment:
    - **Positive**: contains positive words.
    - **Negative**: contains negative words.
    - **Neutral**: contains neither.
- Output: Spark DataFrame with a new **sentiment** column.

### 3.3 Step 3 - PostgreSQL Integration & Analysis

**Objective:** To store processed data in PostgreSQL and run analytical queries.

**Key Actions:**

- Created PostgreSQL table *tweets* with fields:
  - *\_id, author\_id, created\_at, id, text, sentiment*.
- Wrote processed Spark DataFrame to PostgreSQL using JDBC.
- Executed SQL queries to extract insights:
  1. Count of tweets per sentiment: measures overall tone.
  2. Top 10 active users: identifies key contributors.
  3. Sentiment distribution over time: trends by day, hour, and minute.

**Insight:**

- Positive tweets dominated the dataset, indicating a generally optimistic discussion.
- The most active user contributed 3 tweets.

## 4. Results & Findings


- Successfully implemented a real-time social media pipeline.
- Demonstrated seamless integration of NoSQL (MongoDB), big data processing (Spark), and relational databases (PostgreSQL)

- The system can be adapted for:
  - Political opinion tracking.
  - Brand monitoring.
  - Event-driven sentiment changes.

## 5. Challenges & Solutions

Challenge	Solution
Twitter API rate limits	Limited tweet collection to 80 for demonstration.
Large lexicon performance in Spark	Used chunked condition building to avoid performance bottlenecks.
Text noise in tweets	Applied regex-based cleaning to remove unwanted elements.

## 6. Conclusion



This project demonstrated how to build an end-to-end big data pipeline integrating real-time data collection, large-scale processing, and structured storage for analysis. The approach is scalable and adaptable to other domains.

## 7. Future Improvements

- Integrate real-time streaming with Spark Structured Streaming.
- Use machine learning-based sentiment models for higher accuracy.
- Build a dashboard (Power BI / Tableau) for real-time sentiment visualization.
- Expand to multi-language sentiment analysis